

The Hot Path SSA Form in LLVM

Algorithms & Applications


Mohd. Muzzammil¹, Abhay Kumar¹, Sumit Lahiri¹
Awanish Pandey², and Subhajit Roy¹

¹Dept. Of Computer Science & Engineering, IIT Kanpur


²Qualcomm

- 1 Introduction
- 2 The Hot Path SSA (HPSSA) Form
- 3 Constructing the HPSSA Form
- 4 Conclusions


A case for profile-guided optimizations (PGO)



A case for profile-guided optimizations (PGO)




Profile-guided analysis on paths



Summary

- Profile-guided analysis across paths is stronger—can capture correlations between control-flow of basic-blocks
- Collecting path-profiles seems challenging—requires “recording” of a sequence of basic-blocks

Profile-guided analysis on paths



Summary

- Profile-guided analysis across paths is stronger—can capture correlations between control-flow of basic-blocks
- Collecting path-profiles seems challenging—requires “recording” of a sequence of basic-blocks

Ball-Larus Acyclic Profiling [Ball & Larus, MICRO'96]

- Core idea: assign an identifier to each path, that can be calculated efficiently at runtime
- Record frequencies against these identifiers (instead of a sequence of node identifiers)

Ball-Larus Acyclic Profiling [Ball & Larus, MICRO'96]

- Core idea: assign an identifier to each path, that can be calculated efficiently at runtime
- Record frequencies against these identifiers (instead of a sequence of node identifiers)

Capturing still longer paths (k-iteration paths)

- Allows capturing correlations across loop iterations [Roy & Srikant, CGO'09]; a generalization of the Ball-Larus algorithm
- Subsequent work by other groups [D'Elia & Demetrescu, OOPSLA'13]; uses a prefix forest to record BL paths

- **Code understanding**
 - Can expose refactoring opportunities

- **Code understanding**
 - Can expose refactoring opportunities
- **Program testing and verification**
 - Data-driven synthesis of invariants
 - Guided testing for low frequency paths

- **Code understanding**
 - Can expose refactoring opportunities
- **Program testing and verification**
 - Data-driven synthesis of invariants
 - Guided testing for low frequency paths
- **Profile-guided optimizations**

Profile-guided analyses to optimizations

- **Speculation:** *check and retry*
 - eg. value speculation
 - Can impact code size (significant impact w/o speculation support in hardware)

Profile-guided analyses to optimizations

- **Speculation:** *check and retry*
 - eg. value speculation
 - Can impact code size (significant impact w/o speculation support in hardware)
- **Compensation code**
 - eg. superblock scheduling
 - Can impact code size

Profile-guided analyses to optimizations

- **Speculation:** *check and retry*
 - eg. value speculation
 - Can impact code size (significant impact w/o speculation support in hardware)
- **Compensation code**
 - eg. superblock scheduling
 - Can impact code size
- **Error resilient modules**
 - modules for statistical summarization on samples, generate data for ML models
 - No impact on code size

- **Speculation:** *check and retry*
 - eg. value speculation
 - Can impact code size (significant impact w/o speculation support in hardware)
- **Compensation code**
 - eg. superblock scheduling
 - Can impact code size
- **Error resilient modules**
 - modules for statistical summarization on samples, generate data for ML models
 - No impact on code size
- **Optimizations that don't impact correctness**
 - eg. register allocation
 - No impact on code size

- 1 Introduction
- 2 The Hot Path SSA (HPSSA) Form
- 3 Constructing the HPSSA Form
- 4 Conclusions

Why is path-profile-guided analysis hard?

disparate data-structures: program + profile

```
int main()
{
    int a[10] = {{0,2,0}, {1,2,0}, {7,5,0}, {2,5,1}};
    for (int i=0; i<6; i++)
    {
        int max=0;
        int x = a[i][0];
        int y = a[i][1];
        int z = a[i][2];

        int t1 = (x >= y) ? x : y;
        int t2 = (t1 >= z) ? t1 : z;
        int t3 = (z >= x) ? z : x; // bug
        printf("X%d,Y%d,Z%d,X%d,Y%d,Z%d = %d\n", x,y,z,t1,t2,t3, t3);
        if (t1) max = z;
        if (t1 < t2 && t2 < t3) max = z;
        if (t1 < t2 && t2 < t3) max = x;
        if (t2 < t3 && t3 < t1) max = y;
        printf("%d\n", max);
    }
}
```

1-2-4-5	30
21-2-5-5-1	25
1-2-4-5	30
21-2-5-5-1	25
1-2-4-5	30
21-2-5-5-1	25
1-2-4-5	30
21-2-5-5-1	25

Why is path-profile-guided analysis hard?

- There has been enough interest in path-profile-guided analysis and optimizations....
- ...however, designing path-profile-guided variants of traditional optimizations remained hard
- ...hard enough to justify *publications per optimization*
 - Gupta, Benson, Fang. Path profile guided partial dead code elimination using predication. PACT '97.
 - Gupta, Benson, Fang. Path profile guided partial redundancy elimination using speculation. ICCL '98.
 - ...

Can we **weave** profile information into the program representation

Our Objective

Can we **weave** profile information into the program representation
....into a **single, consistent** data-structure

Our Objective

Can we **weave** profile information into the program representation

....into a **single, consistent** data-structure

... that provides the convenience and elegance of an **SSA-like** intermediate form

Can we **weave** profile information into the program representation

....into a **single, consistent** data-structure

... that provides the convenience and elegance of an **SSA-like** intermediate form

...allowing the design of profile-guided versions of “traditional” optimizations with
trivial algorithmic modification of the base algorithms

Can we **weave** profile information into the program representation

....into a **single, consistent** data-structure

... that provides the convenience and elegance of an **SSA-like** intermediate form

...allowing the design of profile-guided versions of “traditional” optimizations with
trivial algorithmic modification of the base algorithms

...providing a **common representation** for both “traditional” as well as profile-guided
analysis and optimizations

The Hot Path SSA Form (HPSSA)

Semantics of a ϕ -function

$$y = \phi(x_1, x_2, \dots, x_n)$$

The Hot Path SSA Form (HPSSA)

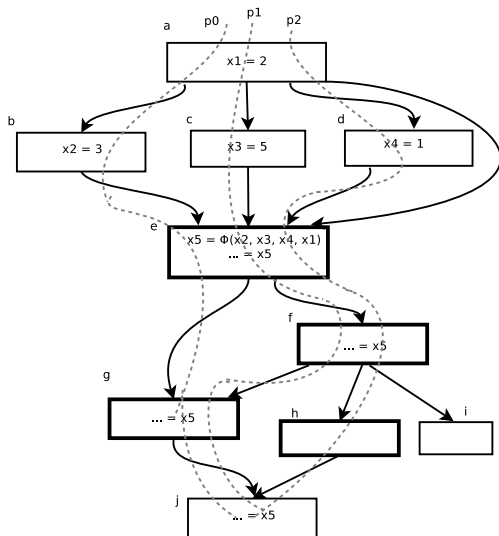
Semantics of a ϕ -function

$$y = \phi(x_1, x_2, \dots, x_n)$$

Semantics of a τ -function

$$\tau(x_0, x_1, x_2, \dots, x_n) = \begin{cases} x_0 & \text{safe interp.} \\ \phi(x_1, x_2, \dots, x_n) & \text{speculative interp.} \end{cases}$$

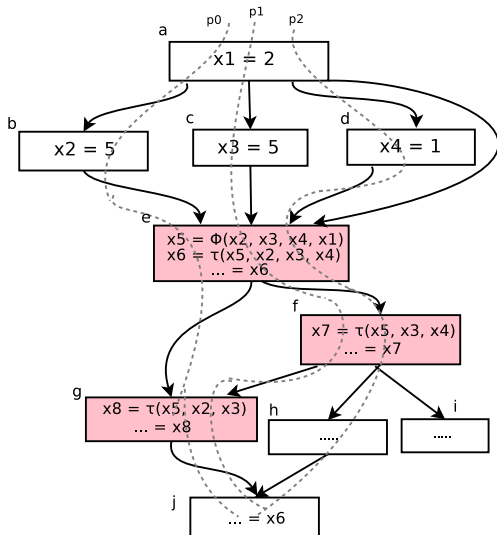
The SSA form



No frequent path carrying:

- def $x_2 = 3$ to use at block **f**
- def $x_4 = 1$ to use at block **g**
- def $x_1 = 2$ to either **f** or **g**

The Hot Path SSA Form



No frequent path carrying:

- def $x_2 = 3$ to use at block **f**
- def $x_4 = 1$ to use at block **g**

Properties

If a program is in the Hot Path SSA form, then,

- each use of a variable is reachable by a single definition; [SSA-like form]

Properties

If a program is in the Hot Path SSA form, then,

- each use of a variable is reachable by a single definition; [SSA-like form]
- **safe interpretation:** [supports traditional analysis]
 - each use of a variable is reachable by the *meet-over-all-paths* reaching definition chains;

The Hot Path SSA Form

Properties

If a program is in the Hot Path SSA form, then,

- each use of a variable is reachable by a single definition; [SSA-like form]
- **safe interpretation:** [supports traditional analysis]
 - each use of a variable is reachable by the *meet-over-all-paths* reaching definition chains;
- **speculative interpretation:** [supports profile-guided analysis]
 - each use of a variable in a basic-block is reachable by the *meet-over-frequent-paths* reaching definitions. ^a

^aor the meet-over-all-paths reaching definition chains, if the use is not reachable from any meet-over-hot-paths reaching definition chain


Speculative Sparse Conditional Constant Propagation (SpecSCCP)

- Introduce new speculative values $\{\dots, 1^s, 2^s, \dots\} \in C^S$
- Operation with *speculative* values result in *speculative* results (with same semantics as base operator)

$$\alpha^s \langle op \rangle \beta = (\alpha \langle op \rangle \beta)^s$$

- Transfer function for τ -functions
($\beta = x_1 \sqcup x_2 \sqcup \dots \sqcup x_n$, i.e. join of speculative args.)

$$\tau(x_0, x_1, \dots, x_n) \sqcup \begin{cases} \top & \text{if } \beta = \top \\ \beta & \text{if } \beta \neq \top \wedge x_0 \sqsubseteq \beta \\ \beta^s & \text{otherwise} \end{cases}$$




Speculative Sparse Conditional Constant Propagation (SpecSCCP)

- Introduce new speculative values $\{\dots, 1^s, 2^s, \dots\} \in C^S$
- Operation with *speculative* values result in *speculative* results (with same semantics as base operator)

$$\alpha^s \langle op \rangle \beta = (\alpha \langle op \rangle \beta)^s$$

- Transfer function for τ -functions
($\beta = x_1 \sqcup x_2 \sqcup \dots \sqcup x_n$, i.e. join of speculative args.)

$$\tau(x_0, x_1, \dots, x_n) \sqcup \begin{cases} \top & \text{if } \beta = \top \\ \beta & \text{if } \beta \neq \top \wedge x_0 \sqsubseteq \beta \\ \beta^s & \text{otherwise} \end{cases}$$



Almost trivial to generate profile-guided variants of standard analyses—an afternoon to “port” SCCP to SpecSCCP!

Speculative Sparse Conditional Constant Propagation (SpecSCCP)

Table 1. Speculative Constants discovered by the SSCP algorithm. (‘~’ indicates *almost*; *grp*, *prg*, & *src* refer to inputs *graphic*, *program* & *source* respectively).

Program	Inpt	Variable Uses		Expression Uses		Total		
		Uses	HitRt	Uses	HitRt	Hits	Misses	HitRt
<i>181.mcf</i>	-	33110	100.00	49665	100.00	82775	0	100.00
<i>175.vpr</i>	-	6938074	100.00	8110837	100.00	15048911	0	100.00
<i>164.gzip</i>	grp	26592	100.00	5	100.00	26597	0	100.00
	prg	17412	100.00	5	100.00	17417	0	100.00
	src	4721	99.98	5	100.00	4725	1	99.98
<i>197.parser</i>	-	165970964	~100.00	340	97.94	165970861	443	~100.00
<i>256.bzip2</i>	grp	132106650	~100.00	938	76.97	132107372	216	~100.00
	prg	100819492	~100.00	6576416	15.67	101849942	5545966	94.84
	src	108134316	~100.00	5256006	17.94	109077366	4312956	96.20

Speculative Sparse Conditional Constant Propagation (SpecSCCP)

Table 1. Speculative Constants discovered by the SSCP algorithm. (‘~’ indicates *almost*; *grp*, *prg*, & *src* refer to inputs *graphic*, *program* & *source* respectively).

Program	Inpt	Variable Uses		Expression Uses		Total		
		Uses	HitRt	Uses	HitRt	Hits	Misses	HitRt
<i>181.mcf</i>	-	33110	100.00	49665	100.00	82775	0	100.00
<i>175.vpr</i>	-	6938074	100.00	8110837	100.00	15048911	0	100.00
<i>164.gzip</i>	<i>grp</i>	26592	100.00	5	100.00	26597	0	100.00
	<i>prg</i>	17412	100.00	5	100.00	17417	0	100.00
	<i>src</i>	4721	99.98	5	100.00	4725	1	99.98
<i>197.parser</i>	-	165970964	~100.00	340	97.94	165970861	443	~100.00
<i>256.bzip2</i>	<i>grp</i>	132106650	~100.00	938	76.97	132107372	216	~100.00
	<i>prg</i>	100819492	~100.00	6576416	15.67	101849942	5545966	94.84
	<i>src</i>	108134316	~100.00	5256006	17.94	109077366	4312956	96.20

Optimized on “train” inputs, results produced on “ref” inputs... Shows the value of profile-guided analysis for code understanding and optimizations...

Profile-Guided Register Allocation [Jain et al., HiPC'16]

<pre># BB#4: addl %edi, 32(%esp) addl \$20, %esi addl \$117, %ebp jmp .LBB0_6 .align 16, 0x90 .LBB0_5: addl \$101, 40(%esp) subl 40(%esp), %ecx addl \$33, %edi addl \$-11, %ebp</pre>	<pre># %if.then # in Loop: Header=B # 4-byte Folded Spill # %if.else # in Loop: Header=B # 4-byte Folded Spill # 4-byte Folded Reload > ></pre>	<pre># BB#4: addl %edi, %ecx addl \$20, %esi addl \$117, %ebp jmp .LBB0_6 .align 16, 0x90 .LBB0_5: addl \$101, 40(%esp) movl 36(%esp), %edx subl 40(%esp), %edx movl %edx, 36(%esp) addl \$33, %edi addl \$-11, %ebp</pre>	<pre># %if.then # in Loop: Header=B # %if.else # in Loop: Header=B # 4-byte Folded Spill # 4-byte Reload # 4-byte Folded Reload # 4-byte Spill</pre>
---	---	---	--

.LBB0_5 is hot: 40(%esp) is cached in %ecx	BB#4 is hot: 32(%esp) is cached in %ecx
--	---

- Implements integer linear programming based profile-guided register allocation
- Base allocator (without profile information) is simplified version of [Goodwin and Wilken, 1996]

Base Allocator

min $\sum_{v \in V, l \in L} S_{v,l} \cdot x_{v,l,\sigma}$, (minimize spills) *such that*

$\forall v \in V \forall l \in L \sum_{r \in R} x_{v,l,r} = 1$ (a variable to at least one location)

$\forall r \in R - \sigma \forall l \in L \sum_{v \in V} x_{v,l,r} \leq 1$ (a register holds at most one variable)

Modelling τ -functions

The cost for $t = \tau(\text{safe}, t_1, t_2, \dots, t_n)$:

$$\text{TauCost} = \sum_{i=1}^{i=n} (\text{loc}(t, t_i) * \text{freq}(t_i))$$

where $\text{loc}(t, t_i)$ is 1 for the locations where t and t_i are different, else 0.

The objective function

The objective function for register allocation over HPSSA is modified to include the *TauCost* discussed above.

The objective function


$$\min \sum_{v \in V, l \in L} S_{v,l} \cdot x_{v,l,\sigma} + \sum_{i=1}^{i=m} \textit{TauCost}_i \quad (1)$$

where, $\textit{TauCost}_i$ is the cost for the i^{th} τ -function (and m is the total number of τ -functions in the program), and


$$S_{v,l} = \begin{cases} 100, & \text{if } v \text{ has further uses along hot path;} \\ 15, & \text{if } v \text{ has further uses only along cold path;} \\ 0, & \text{if } v \text{ has no further use.} \end{cases}$$

Almost a trivial modification to the objective function...


Percentage of spills for different register sizes




(a) gzip




(b) bzip2




(c) mcf



(d) Crafty



(e) GAP



(f) Parser

- 1 Introduction
- 2 The Hot Path SSA (HPSSA) Form
- 3 Constructing the HPSSA Form**
- 4 Conclusions

Definition 1. Hot/Cold Paths

A program path $p : n_1 \rightsquigarrow n_2$ is said to be hot (cold) if the sequence of edges from node n_1 to n_2 appears (does not appear) in any profiled path that occurs frequently in the program profile.

Definition 1. Hot/Cold Paths

A program path $p : n_1 \rightsquigarrow n_2$ is said to be hot (cold) if the sequence of edges from node n_1 to n_2 appears (does not appear) in any profiled path that occurs frequently in the program profile.

Definition 2. Temperature (θ) of a node (edge)

- hot: if the node (edge) is present on a hot path;
- cold: if the node (edge) is not present on any hot path.

Definition 1. Hot/Cold Paths

A program path $p : n_1 \rightsquigarrow n_2$ is said to be hot (cold) if the sequence of edges from node n_1 to n_2 appears (does not appear) in any profiled path that occurs frequently in the program profile.

Definition 2. Temperature (θ) of a node (edge)

- hot: if the node (edge) is present on a hot path;
- cold: if the node (edge) is not present on any hot path.

Definition 3. Hot/Cold Reaching Definitions and Definition Chains

A definition δ at a basic-block n_1 is said to reach a respective use at a basic-block n_2 hot if there exists a hot path from n_1 to n_2 , and δ is not killed along that path. A definition δ at a basic-block n_1 is said to reach a respective use at a basic-block n_2 cold if there does not exist a hot path from n_1 to n_2 , and δ is not killed at least along one cold path from n_1 to n_2 .

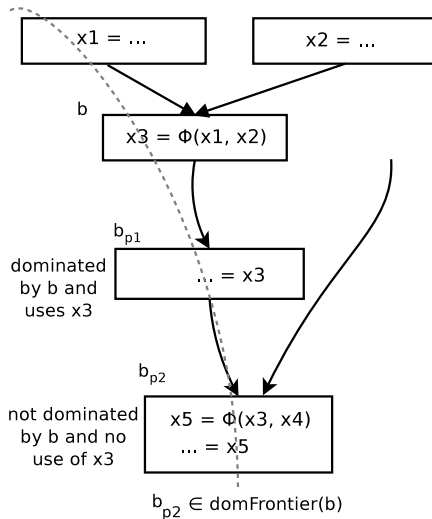
Necessary condition for τ -functions

Lemma 1. A node n requires a τ -function for variable x due to a definition d^x (of a variable x) if

- 1 n is the junction of a hot and a cold path, i.e., paths at different temperatures meet at this node;
- 2 n is reachable by at least two different definitions of the variable x .

Proof. If condition I fails, a τ -function is unnecessary as n can then be reachable by only hot or only cold definitions of x . If condition II fails, a τ -function is again unnecessary as the node is then dominated by a definition of x .

Inserting τ -functions



Definition 4. Thermal Frontier (TF)

For definition d defined at a node u reaching v , v is in the Thermal Frontier of (u,d) , $v \in TF(u, d)$, iff

- 1 the node v is also exposed to a reaching definition d' defined at a node $u \notin Dom(w)$ (w not dominated by u)
- 2 $\theta(u \rightsquigarrow v) \neq \theta(w \rightsquigarrow v)$
- 3 v is the first node on the paths $u \rightsquigarrow v$ and $w \rightsquigarrow v$ that satisfies the above properties.

Definition 4. Thermal Frontier (TF)

For definition d defined at a node u reaching v , v is in the Thermal Frontier of (u,d) , $v \in TF(u, d)$, iff

- 1 the node v is also exposed to a reaching definition d' defined at a node $u \notin Dom(w)$ (w not dominated by u)
- 2 $\theta(u \rightsquigarrow v) \neq \theta(w \rightsquigarrow v)$
- 3 v is the first node on the paths $u \rightsquigarrow v$ and $w \rightsquigarrow v$ that satisfies the above properties.

Theorem

For a set of visible definitions of a variable x at a set of nodes κ , τ -statements are only required at the Iterated Thermal Frontier ITF^x for variable x .

- **Insert ϕ -functions:**
 - insert ϕ -functions at the *iterated dominance frontiers*
- **Insert τ -functions**
 - insert τ -functions at the *iterated thermal frontiers*
- **Allocate arguments to ϕ -functions**
 - use a variable stack to allocate the ϕ -function arguments
- **Allocate arguments to τ -functions**
 - maintain path-sensitive reaching definitions for each program variable corresponding to each hot path on a path-sensitive stack;
 - for each instruction in the program, the algorithm update the respective stack to record the change in the path-sensitive reaching definitions due to the instruction;
 - when a τ -function is encountered, the current set of reaching definitions on the stack is used to allocate the speculative arguments for the τ -function.

Difficulties

- Needs the SSA construction identified, broken into and retrofitted with the HPSSA construction phases.
- The algorithm is quite complex!

HPSSA construction over SSA [Jaiswal et. al., SCOPES'17]

Difficulties

- Needs the SSA construction identified, broken into and retrofitted with the HPSSA construction phases.
- The algorithm is quite complex!

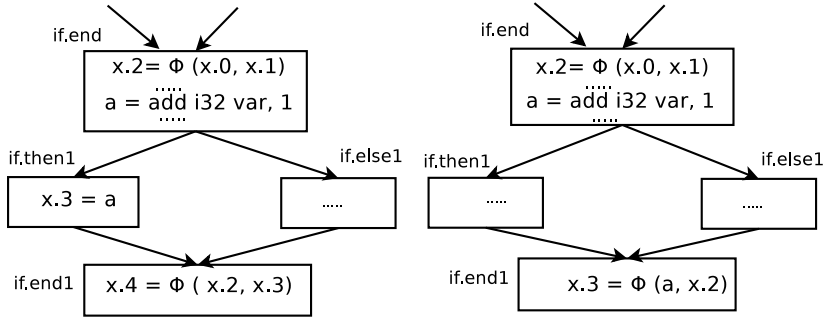
HPSSA over SSA

- **Easily incorporated within existing compilers:** Construction over the SSA form
- **Efficient:** Lesser instructions have to be traversed
- **Simpler:** many constructs are eliminated

A naïve attempt: Mimic [Roy et al.]

- attempt to “recover” the renamed versions of each base variable that is merged by the ϕ -functions;
- then, allocate a single path-sensitive stack for all versions of the same base variable.

Optimized SSA forms

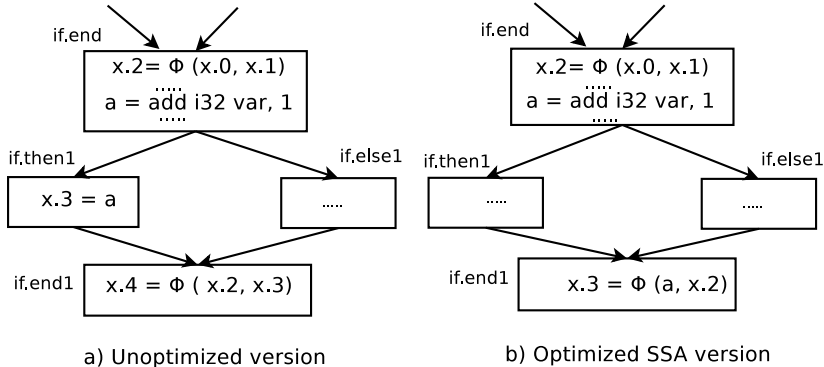


a) Unoptimized version

b) Optimized SSA version

a and x.2 are live simultaneously — hence, cannot be different versions of the same variable

Optimized SSA forms



a and x.2 are live simultaneously — hence, cannot be different versions of the same variable

in the above example, copy propagation breaks the *phi congruence property*...

Shreedhar et al. [SAS'99]

“The occurrences of all resources which belong to the same phi congruence class in a program can be replaced by a representative resource. After the replacement, the phi instruction can be eliminated without violating the semantics of the original program.”

- Sreedhar et al. circumvent the problem by translating the optimized SSA form to the conventional SSA form (that satisfies the phi congruence property) before translating out of SSA.
- **We directly build the HPSSA form over the optimized SSA form!**

- **Insert τ -functions**
 - Insert at Thermal Frontiers
- **Allocate arguments to τ -functions**
 - path-sensitive traversal through the program to identify definitions that reach τ -functions through hot paths
 - constrains its inspection to only the ϕ -functions and the τ -functions

Allocating τ -function arguments

Uses a path-sensitive stack: **phiStack**


- **phiStack** is a stack of **frames**
- each frame $\langle d_i, \xi_i \rangle$ where $\xi_i = \{p_1, p_2, \dots\}$
- support operations:
 - push(frame f, block b)
 - pop(block b)

High-level algorithm

- Load arguments from ϕ - and τ -functions along with their hot path sets on **phiStack**
- Assign the definition from the topmost frame of **phiStack** to any τ -function encountered

- 1 Introduction
- 2 The Hot Path SSA (HPSSA) Form
- 3 Constructing the HPSSA Form
- 4 Conclusions**


Profile-guided analysis on paths



Summary

- Profile-guided analysis across paths is stronger—can capture correlations between control-flow of basic-blocks
- Collecting path-profiles seems challenging—requires “recording” of a sequence of basic-blocks

Profile-guided analyses to optimizations




- **Speculation:** *check and retry*
 - eg. value speculation
 - Can impact code size (significant impact w/o speculation support in hardware)
- **Compensation code**
 - eg. superblock scheduling
 - Can impact code size
- **Error resilient modules**
 - modules for statistical summarization on samples, generate data for ML models
 - No impact on code size
- **Optimizations that don't impact correctness**
 - eg. register allocation
 - No impact on code size

Summary

Introduction The Hot Path SSA (HPSSA) Form Constructing the HPSSA Form Conclusions

Our Objective



Can we **weave** profile information into the program representation

....into a **single, consistent** data-structure

... that provides the convenience and elegance of an **SSA-like** intermediate form

...allowing the design of profile-guided versions of “traditional” optimizations with **trivial algorithmic modification** of the base algorithms

...providing a **common representation** for both “traditional” as well as profile-guided analysis and optimizations

Summary

Introduction
○○○○○○○

The Hot Path SSA (HPSSA) Form
○○○○○○○○●○○○○○

Constructing the HPSSA Form
○○○○○○○○○○○○○○○○

Conclusions
○○○○○○○○○○


Speculative Sparse Conditional Constant Propagation (SSCCP)

- Introduce new speculative values $\{\dots, 1^s, 2^s, \dots\} \in C^s$
- Operation with *speculative* values result in *speculative* results (with same semantics as base operator)
- Transfer function for τ -functions ($\beta = x_1 \sqcup x_2 \sqcup \dots$)

$$\alpha^s \langle op \rangle \beta = (\alpha \langle op \rangle \beta)^s$$

$$\tau(x_0, x_1, \dots, x_n) \sqcup \begin{cases} x_0 \sqcup \top & \text{if } x_0 \sqcup \beta \neq \top \\ \beta & \text{if } x_0 \sqcup \beta = \top \wedge \beta \in C^s \\ \beta^s & \text{otherwise} \end{cases}$$

Almost trivial to generate profile-guided variants of standard analyses—it took us an afternoon to “port” SCCP to SSCCP!




Navigation icons: back, forward, search, etc.

Summary

Introduction 000000 The Hot Path SSA (HPSSA) Form 0000000000000000 Constructing the HPSSA Form 0000000000000000 Conclusions 0000000000

Profile-guided analysis on paths



Summary

- Profile-guided analysis across paths is stronger—can capture correlations between control-flow of basic-blocks
- Collecting path-profiles seems challenging—requires “recording” of a sequence of basic-blocks

Navigation icons

Introduction 000000 The Hot Path SSA (HPSSA) Form 0000000000000000 Constructing the HPSSA Form 0000000000000000 Conclusions 0000000000

Our Objective

Can we **weave** profile information into the program representation

....into a **single, consistent** data-structure

... that provides the convenience and elegance of an **SSA-like** intermediate form

...allowing the design of profile-guided versions of “traditional” optimizations with **trivial algorithmic modification** of the base algorithms

...providing a **common representation** for both “traditional” as well as profile-guided analysis and optimizations

Navigation icons

Introduction 000000 The Hot Path SSA (HPSSA) Form 0000000000000000 Constructing the HPSSA Form 0000000000000000 Conclusions 0000000000

Profile-guided analyses to optimizations

- Speculation:** *check and retry*
 - eg. value speculation
 - Can impact code size (significant impact w/o speculation support in hardware)
- Compensation code**
 - eg. superblock scheduling
 - Can impact code size
- Error resilient modules**
 - modules for statistical summarization on samples, generate data for ML models
 - No impact on code size
- Optimizations that don't impact correctness**
 - eg. register allocation
 - No impact on code size

Navigation icons

Introduction 000000 The Hot Path SSA (HPSSA) Form 0000000000000000 Constructing the HPSSA Form 0000000000000000 Conclusions 0000000000


Speculative Sparse Conditional Propagation (SSCCP)

- Introduce new speculative values $\{\dots, 1^s, 2^s, \dots\} \in C^s$
- Operation with *speculative* values result in *speculative* results (with same semantics as base operator)

$$\alpha^s \langle op \rangle \beta = (\alpha \langle op \rangle \beta)^s$$

- Transfer function for τ -functions ($\beta = x_1 \sqcup x_2 \sqcup \dots$)

$$\tau(x_0, x_1, \dots, x_n) \sqcup \begin{cases} x_0 \sqcup T & \text{if } x_0 \sqcup \beta \neq T \\ \beta & \text{if } x_0 \sqcup \beta = T \wedge \beta \in C^s \\ \beta^s & \text{otherwise} \end{cases}$$

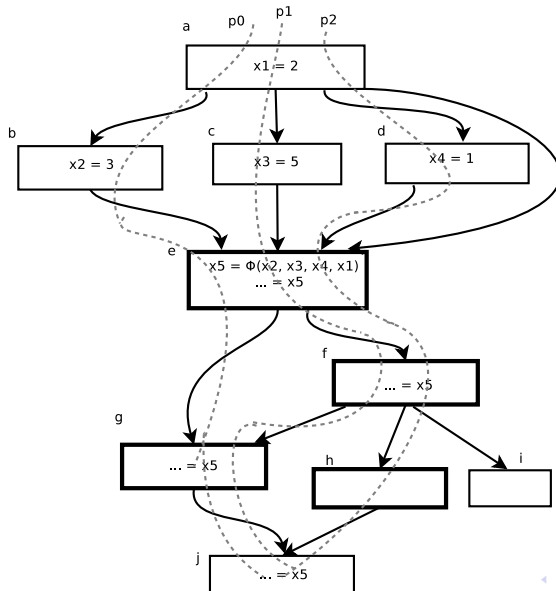


Almost trivial to generate profile-guided variants of standard analyses—it took us an afternoon to “port” SCCP to SSCCP!


Navigation icons

- Subhajit Roy and Y. N. Srikant. Profiling k-Iteration Paths : A Generalization of the Ball-Larus Profiling Algorithm. CGO'09.
- Subhajit Roy and Y. N. Srikant. The Hot Path SSA Form: Extending the Static Single Assignment Form for Speculative Optimizations. CC'10.
- Ramshankar Chouhan, Surender Baswana, Subhajit Roy. Pertinent path profiling: Tracking interactions among relevant statements. CGO'13.
- Era Jain and Subhajit Roy. Phase Directed Compiler Optimizations. HiPC'16.
- Smriti Jaiswal, Praveen Hegde, and Subhajit Roy. Constructing HPSSA over SSA. SCOPES'17.

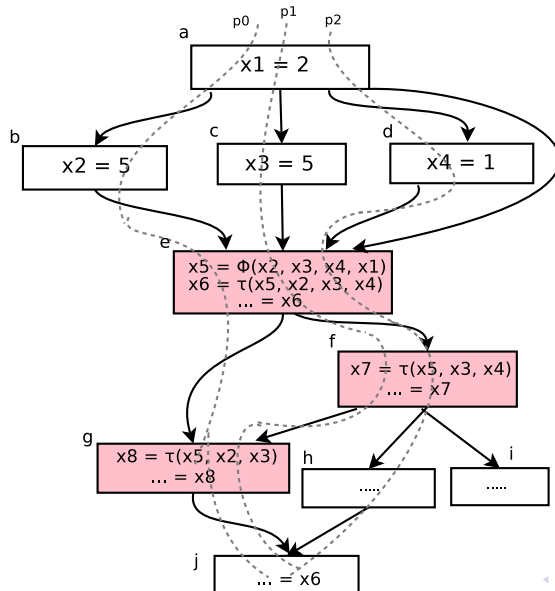
SSA program



τ -functions inserted



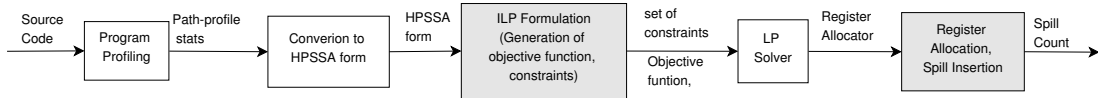
τ -arguments allocated



- We propose a new algorithm that converts SSA programs to HPSSA. Being built over the SSA form, our new algorithm is much more suited for compiler frameworks that offer an SSA-based intermediate representation. This algorithm is also simpler and more efficient than the existing algorithm that builds HPSSA programs from non-SSA programs.
- Our algorithm is capable of operating on optimized SSA forms (resulting from compiler optimizations on the SSA form) that do not satisfy the *Phi Congruence Property*.
- We design an ILP-based path-profile guided register allocation. We provide experimental evaluation of the performance of our approach on eight benchmarks of SPEC CINT2000 benchmark suite.

Questions?

- The occurrences of all resources which belong to the same phi congruence class in a program can be replaced by a representative resource. After the replacement, the phi instruction can be eliminated without violating the semantics of the original program.



```

1: Calculate all caloric connectors in the CFG of the program.
2:  $isInserted(l,b)=\text{false}$  forall instruction  $l$  and basic-block  $b$ 
3: Traverse CFG in topological order, performing the following:
4: for all  $\phi$ -function instruction  $l \in$  basic-block  $b$  do
5:   for all hot path  $p$  passing through  $b$  do
6:     for  $b_p =$  next basic-blocks on  $p$  till  $b_p \notin \text{DomFront}(b)$  do
7:       if  $b_p \in \text{CaloricConnectors} \wedge \neg isInserted(l, b_p)$  then
8:         insert a  $\tau$ -function  $x_{new} = \tau(x_{phi})$  cor-
           responding to the  $\phi$ -function  $l$ :  $x_{phi} =$ 
            $\phi(\dots)$  after all  $\phi$ -function statements
           in block  $b_p$ 
9:          $isInserted(l, b_p) = \text{true}$ 
10:      end if
11:    end for
12:  end for
13: end for

```

```

While processing a basic block  $b$ ,
1: for all  $\phi$ -function  $l$ :  $x_o = \phi(x_1, x_2, \dots, x_n) \in b$  do
2:   for all  $x_i$  for which  $\Omega_{x_i}(b)$  exists do
3:      $\text{phiStack}_{x_i}.\text{push}(\Omega_{x_i}(b), b)$ 
4:   end for
5:   if  $b \in \text{IncubationNode}$  then
6:      $\text{frame} = []$ 
7:     for all  $(x_i, \xi_i) \in \text{phiStack}_{x_i}(b).Top$  do
8:        $\text{frame} = \text{frame.add}(x_i, \xi_i \cup \text{IncubPaths}(b))$ 
9:     end for
10:    for all arguments  $x_i$  in  $\phi$ -function  $l$  do
11:      if  $x_i$  is a concrete defn.  $\wedge b_p \rightarrow b$  an incoming edge then
12:         $\text{frame} = \text{frame.add}(x_i, \text{pathSet}(b_p \rightarrow b) \cup \text{IncubPaths}(b))$ 
13:      end if
14:      if  $x_i$  is a  $\phi$ -argument corresponding to a hot backedge then
15:         $\text{frame} = \text{frame.add}(x_i, \text{IncubPaths}(b))$ 
16:      end if
17:    end for
18:     $\text{phiStack}_x.\text{push}(\text{frame}, b)$ 
19:  end if
20: end for
21: for all  $\tau$ -function instructions  $l \in b$  do
22:   for all  $(x_i, \xi_i) \in \text{phiStack}_{x_i}(b).Top$  do
23:    if  $\xi_i \cap \text{pathSet}(b) \neq \emptyset$  then
24:      add  $(x_i, \xi_i \cap \text{pathSet}(b))$  to speculative arguments of  $l$ 
25:    end if
26:  end for
27: end for
28: for all outgoing edge  $b \rightarrow b_s$  do
29:   if  $b_s$  is a join node then
30:    for all  $(x_i, \xi_i) \in \text{phiStack}_{x_i}.Top$  do
31:     if  $\xi_i \cap \text{pathSet}(b \rightarrow b_s) \neq \emptyset$  then
32:        $\Omega_{x_i}(b_s) = \xi_i \cap \text{pathSet}(b \rightarrow b_s)$ 
33:     end if
34:   end for
35: end if
36: end for
37: Make a recursive call on the children of node  $b$  in the dominator tree, traversing
   them in topological order of the nodes in the control-flow graph order.
38: Pop off all frames pushed by block  $b$  from all  $\text{phiStacks}$ 

```